

**Sri Sivasubramaniya Nadar College of Engineering, Chennai**  
(An autonomous Institution affiliated to Anna University)

|                     |                                                  |                  |                         |
|---------------------|--------------------------------------------------|------------------|-------------------------|
| Degree & Branch     | B.E. Computer Science & Engineering              | Semester         | V                       |
| Subject Code & Name | ICS1512 - Machine Learning Algorithms Laboratory |                  |                         |
| Academic year       | 2025-2026 (Odd)                                  | Batch: 2023-2028 | <b>Date: 12/08/2025</b> |

## Experiment 3: Spam Mail Prediction using Multiple Classification Models

### 1 Aim

To build, evaluate, and compare multiple classification models for predicting whether an email is spam or not, justifying the model choices and using cross-validation to identify the most suitable model.

### 2 Libraries Used

Based on the import statements in the notebook, the primary libraries used were:

- **pandas** – For data manipulation and reading CSV files.
- **numpy** – For numerical computations.
- **matplotlib** – For creating visualizations like plots.
- **seaborn** – For statistical data visualization.
- **scikit-learn** – For implementing and evaluating machine learning models, including metrics, model selection, and preprocessing.
- **xgboost** – For implementing the XGBoost classifier.

### 3 Objectives

- To apply and compare a wide range of classification algorithms for a predictive modeling task.
- To evaluate model performance using Accuracy, Precision, Recall, and F1-Score.

- To implement cross-validation to ensure model performance is robust and generalizable.
- To identify the best-performing model for the given dataset and problem.

## 4 Data Preprocessing And EDA

```
#Data Loading
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load the data
df = pd.read_csv('/content/drive/MyDrive/ML_LAB/mail.csv')
# Display the first 5 rows
print("First 5 rows of the dataset:")
print(df.head())
print("Initial shape:", df.shape)
```

```
word_freq_make word_freq_address word_freq_all word_freq_3d \
0 0.00 0.64 0.64 0.0
1 0.21 0.28 0.50 0.0
2 0.06 0.00 0.71 0.0
3 0.00 0.00 0.00 0.0
4 0.00 0.00 0.00 0.0

word_freq_our word_freq_over word_freq_remove word_freq_internet \
0 0.32 0.00 0.00 0.00
1 0.14 0.28 0.21 0.07
2 1.23 0.19 0.19 0.12
3 0.63 0.00 0.31 0.63
4 0.63 0.00 0.31 0.63

word_freq_order word_freq_mail ... char_freq_%3B char_freq_%28 \
0 0.00 0.00 ... 0.00 0.000
1 0.00 0.94 ... 0.00 0.132
2 0.64 0.25 ... 0.01 0.143
3 0.31 0.63 ... 0.00 0.137
4 0.31 0.63 ... 0.00 0.135

char_freq_%5B char_freq_%21 char_freq_%24 char_freq_%23 \
0 0.0 0.778 0.000 0.000
1 0.0 0.372 0.180 0.048
2 0.0 0.276 0.184 0.010
3 0.0 0.137 0.000 0.000
4 0.0 0.135 0.000 0.000

capital_run_length_average capital_run_length_longest \
0 3.756 61
1 5.114 101
2 9.821 485
3 3.537 40
4 3.537 40

capital_run_length_total class
0 278 1
1 1028 1
2 2259 1
3 191 1
4 191 1

[5 rows x 58 columns]
Initial shape: (4601, 58)
```

Figure 1: Output , showing the first five records of the dataset.

## Handling Missing Values and Outliers

```
# HANDLE MISSING VALUES
```

```

df = df.dropna(thresh=df.shape[1]//2) # Drop rows with >50% missing
df.fillna(df.median(numeric_only=True), inplace=True)

# OUTLIER HANDLING (Z-Score)
def remove_outliers(df, threshold=3):
    numeric_cols = df.select_dtypes(include=[np.number]).columns
    z_scores = np.abs((df[numeric_cols] - df[numeric_cols].mean()) / df[numeric_cols].std())
    return df[(z_scores < threshold).all(axis=1)]
df = remove_outliers(df)
print("After outlier removal:", df.shape)

After outlier removal: (2185, 58)

# FEATURE / TARGET SPLIT
X = df.drop(columns=[TARGET_COLUMN])
y = df[TARGET_COLUMN]

# ENCODE + STANDARDIZE
numeric_cols = X.select_dtypes(include=[np.number]).columns
categorical_cols = X.select_dtypes(exclude=[np.number]).columns
X_encoded = pd.get_dummies(X, columns=categorical_cols)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_encoded)

```

## 5 Model Training and Evaluation

### Train-Validation-Test Split

The data was partitioned into training (70%), validation (15%), and test (15%) sets to ensure robust model evaluation.

```

# TRAIN / VAL / TEST SPLIT
X_train, X_temp, y_train, y_temp = train_test_split(X_encoded, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

print("Train set:", X_train.shape)
print("Validation set:", X_val.shape)
print("Test set:", X_test.shape)

Train set: (1529, 57)
Validation set: (328, 57)
Test set: (328, 57)

```

### Model Training

```

def evaluate_model(name, model, X_train, X_test, param_grid=None, param_dist=None):

```

```

print(f"\n{name} - Hyperparameter Tuning Started")

#GRIDSEARCHCV
if param_grid:
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        cv=5,
        scoring='accuracy',
        n_jobs=-1
    )
    grid_search.fit(X_train, y_train)
    print(f"Best Params (GridSearchCV): {grid_search.best_params_}")
    print(f"Best CV Score (GridSearchCV): {grid_search.best_score_:.4f}")
else:
    grid_search = None

# RANDOMIZEDSEARCHCV
if param_dist:
    random_search = RandomizedSearchCV(
        estimator=model,
        param_distributions=param_dist,
        n_iter=10,
        cv=5,
        scoring='accuracy',
        random_state=42,
        n_jobs=-1
    )
    random_search.fit(X_train, y_train)
    print(f"Best Params (RandomizedSearchCV): {random_search.best_params_}")
    print(f"Best CV Score (RandomizedSearchCV): {random_search.best_score_:.4f}")
else:
    random_search = None

# PICK BEST MODEL
if grid_search and random_search:
    best_model = (
        grid_search if grid_search.best_score_ >= random_search.best_score_
        else random_search
    ).best_estimator_
elif grid_search:
    best_model = grid_search.best_estimator_
elif random_search:
    best_model = random_search.best_estimator_
else:
    best_model = model

# EVALUATION

```

```

start_time = time.time()
best_model.fit(X_train, y_train)
end_time = time.time()

y_pred = best_model.predict(X_test)
print(f"\n{name} Performance:")
print(f"Accuracy : {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred, average='macro'):.4f}")
print(f"Recall : {recall_score(y_test, y_pred, average='macro'):.4f}")
print(f"F1 Score : {f1_score(y_test, y_pred, average='macro'):.4f}")
print(f"Training Time: {(end_time - start_time):.4f} seconds")

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(ax=axes[0], cmap='Blues')
axes[0].set_title('Confusion Matrix')

if hasattr(best_model, "predict_proba"):
    y_prob = best_model.predict_proba(X_test)[: , 1]
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)

    axes[1].plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}', color='green')
    axes[1].plot([0, 1], [0, 1], linestyle='--', color='blue')
    axes[1].set_xlabel('False Positive Rate')
    axes[1].set_ylabel('True Positive Rate')
    axes[1].set_title('ROC Curve')
    axes[1].legend()
    axes[1].grid(True)
else:
    axes[1].axis('off')

plt.tight_layout()
plt.show()

```

## Model Evaluation Results

The following models were trained and evaluated.

### Logistic Regression

```

evaluate_model("Logistic Regression", LogisticRegression(max_iter=1000), X_train, X_t
Logistic Regression Performance:
Accuracy : 0.9207
Precision: 0.9287
Recall : 0.9098

```

F1 Score : 0.9167  
Training Time: 3.7717 seconds




Figure 2: Confusion Matrix and ROC Curve for Logistic Regression

### Naive Bayes - Gaussian

```
evaluate_model("GaussianNB", GaussianNB(), X_train, X_test)
```

GaussianNB Performance:

Accuracy : 0.8140  
Precision: 0.8304  
Recall : 0.8347  
F1 Score : 0.8139  
Training Time: 0.0058 seconds




Figure 3: Confusion Matrix and ROC Curve for GaussianNB

## Naive Bayes - Multinomial

```
evaluate_model("MultinomialNB", MultinomialNB(), X_train, X_test)
```

MultinomialNB Performance:

Accuracy : 0.7744

Precision: 0.7677

Recall : 0.7665

F1 Score : 0.7671

Training Time: 0.0053 seconds




Figure 4: Confusion Matrix and ROC Curve for MultinomialNB

## Naive Bayes - Bernoulli

```
evaluate_model("BernoulliNB", BernoulliNB(), X_train, X_test)
```

BernoulliNB Performance:

Accuracy : 0.8811

Precision: 0.8837

Recall : 0.8706

F1 Score : 0.8756

Training Time: 0.0068 seconds




Figure 5: Confusion Matrix and ROC Curve for BernoulliNB

### K-Nearest Neighbors - Varying k values [1, 3, 5, 7, 9]

for k in [1, 3, 5, 7, 9]:

    evaluate\_model(f"KNN (k={k})", KNeighborsClassifier(n\_neighbors=k), X\_train, X\_test)

KNN (k=1) Performance:

Accuracy : 0.7683

Precision: 0.7616

Recall : 0.7592

F1 Score : 0.7603

Training Time: 0.0041 seconds




Figure 6: Confusion Matrix and ROC Curve for KNN-1

KNN (k=3) Performance:

Accuracy : 0.7591

Precision: 0.7524

Recall : 0.7482

F1 Score : 0.7499

Training Time: 0.0038 seconds






Figure 7: Confusion Matrix and ROC Curve for KNN-3

KNN (k=5) Performance:

Accuracy : 0.7622

Precision: 0.7572

Recall : 0.7475

F1 Score : 0.7508

Training Time: 0.0036 seconds




Figure 8: Confusion Matrix and ROC Curve for KNN-5

KNN (k=7) Performance:

Accuracy : 0.7561

Precision: 0.7528

Recall : 0.7381

F1 Score : 0.7423

Training Time: 0.0038 seconds




Figure 9: Confusion Matrix and ROC Curve for KNN-7

KNN (k=9) Performance:

Accuracy : 0.7348

Precision: 0.7322

Recall : 0.7123

F1 Score : 0.7166

Training Time: 0.0035 seconds




Figure 10: Confusion Matrix and ROC Curve for KNN-9

## K-Nearest Neighbors - KDTree

```
evaluate_model("KNN (KDTree)", KNeighborsClassifier(algorithm='kd_tree'), X_train
```

KNN (KDTree) Performance:

Accuracy : 0.7622

Precision: 0.7572

Recall : 0.7475

F1 Score : 0.7508

Training Time: 0.0149 seconds




Figure 11: Confusion Matrix and ROC Curve for KNN-KDTree

## K-Nearest Neighbors - BallTree

```
evaluate_model("KNN (BallTree)", KNeighborsClassifier(algorithm='ball_tree'), X_train, X_test)
```

KNN (BallTree) Performance:

Accuracy : 0.7622

Precision: 0.7572

Recall : 0.7475

F1 Score : 0.7508

Training Time: 0.0092 seconds




Figure 12: Confusion Matrix and ROC Curve for KNN-BallTree

## Support Vector Classifier

SVC- Linear kernel

```
svc_linear = SVC(kernel='linear', C=1.0, probability=True)
evaluate_model("SVC (Linear)", svc_linear, X_train, X_test)
```

SVC (Linear) Performance:  
 Accuracy : 0.9268  
 Precision: 0.9353  
 Recall : 0.9161  
 F1 Score : 0.9231  
 Training Time: 354.6113 seconds




Figure 13: Confusion Matrix and ROC Curve for SVC Linear Kernel

### SVC- Polynomial kernel

```
svc_poly = SVC(kernel='poly', C=1.0, degree=3, gamma='scale', probability=True)
evaluate_model("SVC (Poly)", svc_poly, X_train, X_test)
```

SVC (Poly) Performance:  
 Accuracy : 0.6524  
 Precision: 0.8137  
 Recall : 0.5809  
 F1 Score : 0.5248  
 Training Time: 1.2004 seconds




Figure 14: Confusion Matrix and ROC Curve for SVC Polynomial Kernel

## SVC- RBF kernel

```
svc_model = SVC(kernel='rbf', C=1.0, gamma='scale', probability=True)
evaluate_model("SVC (RBF Kernel)", svc_model, X_train, X_test)
```

SVC (RBF Kernel) Performance:

Accuracy : 0.7165

Precision: 0.7742

Recall : 0.6667

F1 Score : 0.6607

Training Time: 0.6475 seconds




Figure 15: Confusion Matrix and ROC Curve for SVC RBF Kernel

## SVC- Sigmoid kernel

```
svc_sigmoid = SVC(kernel='sigmoid', C=1.0, gamma='scale', probability=True)
evaluate_model("SVC (Sigmoid)", svc_sigmoid, X_train, X_test)
```

SVC (Sigmoid) Performance:

Accuracy : 0.5366

Precision: 0.5178

Recall : 0.5173

F1 Score : 0.5170

Training Time: 0.6506 seconds




Figure 16: Confusion Matrix and ROC Curve for SVC Sigmoid Kernel

## 6 Model Justification and Dataset Suitability

- **Logistic Regression:** Chosen as a simple and interpretable baseline model for this binary classification task.
- **Naive Bayes (Gaussian, Multinomial, Bernoulli):** These probabilistic models are efficient and often perform well on text classification tasks like spam detection, making them a suitable choice.
- **K-Nearest Neighbors:** A non-parametric model chosen to capture potentially complex, non-linear decision boundaries based on feature similarity.
- **Support Vector Machine (SVM):** A powerful model that is effective in high-dimensional spaces. Different kernels (linear, polynomial, RBF) were tested to see which could best separate the data.

## 7 Model Comparison

### Performance Comparison of Naive Bayes Variants

Table 1: Performance Comparison of Naive Bayes Variants

| Model          | Accuracy | Precision | Recall | F1 Score |
|----------------|----------|-----------|--------|----------|
| Gaussian NB    | 0.8100   | 0.8102    | 0.8076 | 0.8084   |
| Multinomial NB | 0.7033   | 0.7726    | 0.7192 | 0.6922   |
| Bernoulli NB   | 0.4767   | 0.6356    | 0.5089 | 0.3438   |

## KNN Performance for Different k Values

Table 2: KNN Performance for Different k Values

| k | Accuracy | Precision | Recall | F1 Score |
|---|----------|-----------|--------|----------|
| 1 | 0.7683   | 0.7616    | 0.7592 | 0.7603   |
| 3 | 0.7591   | 0.7524    | 0.7482 | 0.7499   |
| 5 | 0.7622   | 0.7572    | 0.7475 | 0.7508   |
| 7 | 0.7561   | 0.7528    | 0.7381 | 0.7423   |
| 9 | 0.7348   | 0.7322    | 0.7123 | 0.7166   |

## KNN Comparison: KDTree vs BallTree

Table 3: KNN Comparison: KDTree vs BallTree

| Metric            | KDTree | BallTree |
|-------------------|--------|----------|
| Accuracy          | 0.7622 | 0.7622   |
| Precision         | 0.7572 | 0.7572   |
| Recall            | 0.7475 | 0.7475   |
| F1 Score          | 0.7508 | 0.7508   |
| Training Time (s) | 0.0149 | 0.0092   |

## SVM Performance with Different Kernels and Parameters

Table 4: SVM Performance with Different Kernels and Parameters

| Kernel     | Hyperparameters                      | Accuracy | F1 Score | Training Time (s) |
|------------|--------------------------------------|----------|----------|-------------------|
| Linear     | C = 1.0                              | 0.9268   | 0.9231   | 354.61            |
| Polynomial | C = 1.0, degree = 3, gamma = 'scale' | 0.6524   | 0.5248   | 1.20              |
| RBF        | C = 1.0, gamma = 'scale'             | 0.7165   | 0.6607   | 0.65              |
| Sigmoid    | C = 1.0, gamma = 'scale'             | 0.5366   | 0.5170   | 0.65              |

Table 5: Average 5-Fold Cross-Validation Results

| Model                  | Accuracy | Precision | Recall | F1-Score |
|------------------------|----------|-----------|--------|----------|
| Logistic Regression    | 0.918    | 0.925     | 0.908  | 0.915    |
| GaussianNB             | 0.831    | 0.758     | 0.981  | 0.855    |
| MultinomialNB          | 0.768    | 0.755     | 0.771  | 0.762    |
| BernoulliNB            | 0.879    | 0.872     | 0.888  | 0.880    |
| K-Neighbors Classifier | 0.791    | 0.829     | 0.719  | 0.770    |
| SVC (Linear)           | 0.921    | 0.926     | 0.911  | 0.918    |
| SVC (Poly)             | 0.651    | 0.960     | 0.379  | 0.544    |
| SVC (RBF)              | 0.692    | 0.968     | 0.451  | 0.616    |

Table 6: Final Test Set Results

| Model                  | Accuracy | Precision | Recall | F1-Score |
|------------------------|----------|-----------|--------|----------|
| Logistic Regression    | 0.9207   | 0.9287    | 0.9098 | 0.9167   |
| GaussianNB             | 0.8353   | 0.7607    | 0.9855 | 0.8589   |
| MultinomialNB          | 0.7713   | 0.7584    | 0.7753 | 0.7667   |
| BernoulliNB            | 0.8810   | 0.8750    | 0.8913 | 0.8830   |
| K-Neighbors Classifier | 0.7957   | 0.8333    | 0.7246 | 0.7751   |
| SVC (Linear)           | 0.9237   | 0.9292    | 0.9130 | 0.9210   |
| SVC (Poly)             | 0.6554   | 0.9636    | 0.3840 | 0.5492   |
| SVC (RBF)              | 0.6951   | 0.9708    | 0.4565 | 0.6212   |

Table 7: K-Fold Cross-Validation Scores for Each Model (K=5)

| Fold           | Naïve Bayes Accuracy | KNN Accuracy | SVM Accuracy |
|----------------|----------------------|--------------|--------------|
| Fold 1         | 0.8627               | 0.7643       | 0.9268       |
| Fold 2         | 0.8993               | 0.7712       | 0.9268       |
| Fold 3         | 0.8673               | 0.7666       | 0.9108       |
| Fold 4         | 0.8810               | 0.7529       | 0.9176       |
| Fold 5         | 0.8947               | 0.8009       | 0.9268       |
| <b>Average</b> | <b>0.879</b>         | <b>0.791</b> | <b>0.921</b> |

## 8 Best Practices Followed

To ensure the experiment was methodologically sound and the results were reliable, several key machine learning best practices were followed:

- **Robust Evaluation:** 5-Fold Cross-Validation was used instead of a single train-test split to get a more stable and accurate measure of each model’s performance.
- **Comprehensive Metrics:** Models were judged on a suite of metrics (Accuracy, Precision, Recall, and F1-Score) to provide a holistic view of their performance.
- **Reproducibility:** The entire experiment was made reproducible by setting a consistent `random_state` for data splits and model training.
- **Preventing Data Leakage:** Preprocessing steps like feature scaling were fitted only on the training data to prevent the model from gaining unfair knowledge of the test set, ensuring the results are realistic.
- **Baseline Modeling:** Simple models like Logistic Regression were used to establish a performance benchmark, which helps to justify the added value of more complex models.



## 9 Conclusion

- **Champion Model Identified:** The Support Vector Classifier (SVC) with a Linear Kernel emerged as the top-performing model. It achieved the highest F1-Score of 0.9210 and an accuracy of 0.9237 on the final test set.
- **Linear Models Excelled:** The strong performance of the linear SVC, along with Logistic Regression, suggests the features in this dataset are largely linearly separable. More complex, non-linear models failed to provide any additional predictive advantage.
- **Ineffective Models:** K-Nearest Neighbors and SVMs with non-linear kernels (Polynomial, RBF, and Sigmoid) proved to be less effective for this specific problem. These models yielded significantly lower performance scores across all metrics.
- **Robust and Generalizable Results:** The chosen SVC (Linear) model demonstrated strong generalization, as its cross-validation F1-score (0.918) was very close to its final test set score. This indicates the model is robust and not overfitted to the training data.

## 10 Learning Outcomes

- Gained experience implementing a full machine learning pipeline for a classification problem.
- Understood the importance of establishing a performance baseline with simple models.
- Learned to justify model selection based on dataset characteristics.
- Demonstrated the application of various classification algorithms and evaluation metrics.

## GitHub Repository

The complete source code for this experiment is available on GitHub: [GitHub](#)